

Neural Networks Hello World + Assignments 2, 3

(Neural Networks Implementation and Application Tutorial)

Vilém Zouhar, Noon Pokaratsiri Goldstein

24th, 25th November 2021

Overview

- Assignment 2
- Gradient
- PyTorch's Autograd
- NN Hello World
- Assignment 3

Assignment 2

TODO

Optimization

Gradient 🤔

- What is it?

Optimization

Gradient 🤔

- What is it?
- How do we denote it?

Optimization

Gradient 🤔

- What is it?
- How do we denote it?
 - ▶ $\nabla f(p) = [\frac{\delta f}{\delta x_1}(p), \dots, \frac{\delta f}{\delta x_k}(p)]$

Optimization

Gradient 🤔

- What is it?
- How do we denote it?
 - ▶ $\nabla f(p) = [\frac{\delta f}{\delta x_1}(p), \dots, \frac{\delta f}{\delta x_k}(p)]$
- Why is it important?

Optimization

Gradient 🤔

- What is it?
- How do we denote it?
 - ▶ $\nabla f(p) = [\frac{\delta f}{\delta x_1}(p), \dots, \frac{\delta f}{\delta x_k}(p)]$
- Why is it important?
 - ▶ Optimization

Optimization

Gradient 🤔

- What is it?
- How do we denote it?
 - ▶ $\nabla f(p) = [\frac{\delta f}{\delta x_1}(p), \dots, \frac{\delta f}{\delta x_k}(p)]$
- Why is it important?
 - ▶ Optimization

Optimization

Gradient 🤔

- What is it?
- How do we denote it?
 - ▶ $\nabla f(p) = [\frac{\delta f}{\delta x_1}(p), \dots, \frac{\delta f}{\delta x_k}(p)]$
- Why is it important?
 - ▶ Optimization




Figure 1: Function parameter landscape from [1]

Optimization

Few questions 🧐

- How does step/gradient-based optimization work?

Optimization

Few questions 🤔

- How does step/gradient-based optimization work?
- How is the step size determined?

Optimization

Few questions 🤔

- How does step/gradient-based optimization work?
- How is the step size determined?
- Why do we subtract the gradient and not add it?

Optimization

Few questions 🤔

- How does step/gradient-based optimization work?
- How is the step size determined?
- Why do we subtract the gradient and not add it?
- If we start in different places will we always find the same spot?

Optimization

Few questions 🤔

- How does step/gradient-based optimization work?
- How is the step size determined?
- Why do we subtract the gradient and not add it?
- If we start in different places will we always find the same spot?
- Will we always find the global minimum?

Optimization

Few questions 🤔

- How does step/gradient-based optimization work?
- How is the step size determined?
- Why do we subtract the gradient and not add it?
- If we start in different places will we always find the same spot?
- Will we always find the global minimum?

Optimization

Few questions 🧐

- How does step/gradient-based optimization work?
- How is the step size determined?
- Why do we subtract the gradient and not add it?
- If we start in different places will we always find the same spot?
- Will we always find the global minimum?



Figure 2: Function parameter landscape from [2]

Autograd & PyTorch

TODO

Assignment 3

- Any questions?

Resources

- [1] Optimization & landscapes offconvex.org/2018/11/07/optimization-beyond-landscape/
- [2] Optimization Introduction by Scipy
scipy-lectures.org/advanced/mathematical_optimization/