

Neural Networks Hello World + Assignments 2, 3

(Neural Networks Implementation and Application Tutorial)

Vilém Zouhar, Noon Pokaratsiri Goldstein

24th, 25th November 2021

Overview

- Assignment 2
- Gradient
- PyTorch's Autograd
- NN Hello World
- Assignment 3

Assignment 2

TODO

Optimization

Gradient 🤔

- What is it?

Optimization

Gradient 🤔

- What is it?
- How do we denote it?

Optimization

Gradient 🤔

- What is it?
- How do we denote it?
 - ▶ $\nabla f(p) = [\frac{\delta f}{\delta x_1}(p), \dots, \frac{\delta f}{\delta x_k}(p)]$

Optimization

Gradient 🤔

- What is it?
- How do we denote it?
 - ▶ $\nabla f(p) = [\frac{\delta f}{\delta x_1}(p), \dots, \frac{\delta f}{\delta x_k}(p)]$
- Why is it important?

Optimization

Gradient 🤔

- What is it?
- How do we denote it?
 - ▶ $\nabla f(p) = [\frac{\delta f}{\delta x_1}(p), \dots, \frac{\delta f}{\delta x_k}(p)]$
- Why is it important?
 - ▶ Optimization

Optimization

Gradient 🤔

- What is it?
- How do we denote it?
 - ▶ $\nabla f(p) = [\frac{\delta f}{\delta x_1}(p), \dots, \frac{\delta f}{\delta x_k}(p)]$
- Why is it important?
 - ▶ Optimization

Optimization

Gradient 🤔

- What is it?
- How do we denote it?
 - ▶ $\nabla f(p) = [\frac{\delta f}{\delta x_1}(p), \dots, \frac{\delta f}{\delta x_k}(p)]$
- Why is it important?
 - ▶ Optimization

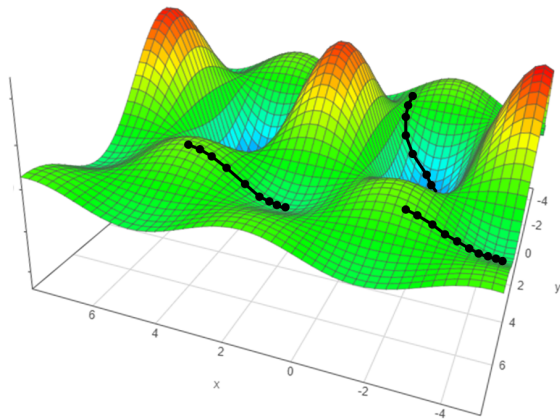


Figure 1: Function parameter landscape from [1]

Optimization

Few questions 🧐

- How does step/gradient-based optimization work?

Optimization

Few questions 🤔

- How does step/gradient-based optimization work?
- How is the step size determined?

Optimization

Few questions 🤔

- How does step/gradient-based optimization work?
- How is the step size determined?
- Why do we subtract the gradient and not add it?

Optimization

Few questions 🤔

- How does step/gradient-based optimization work?
- How is the step size determined?
- Why do we subtract the gradient and not add it?
- If we start in different places will we always find the same spot?

Optimization

Few questions 🤔

- How does step/gradient-based optimization work?
- How is the step size determined?
- Why do we subtract the gradient and not add it?
- If we start in different places will we always find the same spot?
- Will we always find the global minimum?

Optimization

Few questions 🤔

- How does step/gradient-based optimization work?
- How is the step size determined?
- Why do we subtract the gradient and not add it?
- If we start in different places will we always find the same spot?
- Will we always find the global minimum?

Optimization

Few questions 🧐

- How does step/gradient-based optimization work?
- How is the step size determined?
- Why do we subtract the gradient and not add it?
- If we start in different places will we always find the same spot?
- Will we always find the global minimum?



Figure 2: Function parameter landscape from [2]

Autograd & PyTorch

How to get the gradient at $(x, y) = (2, 3)$ of $x \cdot y + \sin(\pi \cdot x)$?

- By hand 🤔

Autograd & PyTorch

How to get the gradient at $(x, y) = (2, 3)$ of $x \cdot y + \sin(\pi \cdot x)$?

- By hand 😓
- Autograd 🌟

Autograd & PyTorch

How to get the gradient at $(x, y) = (2, 3)$ of $x \cdot y + \sin(\pi \cdot x)$?

- By hand 😓
- Autograd 🌟

Autograd & PyTorch

How to get the gradient at $(x, y) = (2, 3)$ of $x \cdot y + \sin(\pi \cdot x)$?

- By hand 😓
- Autograd 🌟🌟

By hand

$$\frac{\delta}{\delta x} = y + \pi \cdot \cos(\pi \cdot x) \rightarrow 3 + \pi$$

$$\frac{\delta}{\delta y} = x \rightarrow 2$$

$$\nabla f(2, 3) \rightarrow (3 + \pi, 2)$$

Autograd & PyTorch

How to get the gradient at $(x, y) = (2, 3)$ of $x \cdot y + \sin(\pi \cdot x)$?

- By hand 🤔
- Autograd 🌟😊

By hand

$$\frac{\delta}{\delta x} = y + \pi \cdot \cos(\pi \cdot x) \rightarrow 3 + \pi$$

$$\frac{\delta}{\delta y} = x \rightarrow 2$$

$$\nabla f(2, 3) \rightarrow (3 + \pi, 2)$$

Autograd

```
import torch
import numpy as np
x = torch.tensor(2.0, requires_grad=True)
y = torch.tensor(3.0, requires_grad=True)
out = x*y + torch.sin(np.pi*x)
out.backward() # trigger gradient computation
assert np.isclose(x.grad, 3+np.pi)
assert np.isclose(y.grad, 2)
```

Assignment 3

- Any questions?

Resources

- [1] Optimization & landscapes offconvex.org/2018/11/07/optimization-beyond-landscape/
- [2] Optimization Introduction by Scipy
scipy-lectures.org/advanced/mathematical_optimization/