

# Batching & Optimizers + Assignments 6, 7

(Neural Networks Implementation and Application Tutorial)

Vilém Zouhar, Noon Pokaratsiri Goldstein

4th, 5th January 2022

# Overview

- Assignment 6
- Batching
- Optimization algorithms
- Assignment 7

# Assignment 6

- What was the hardest part?

# Batching

## Gradient descent 🧐

- What is it?
- What do we base the gradient estimates on?

## What's the difference? 🤔

- Batch, stochastic, mini-batch gradient descent

## What's the batch size for dataset size $d$ ? 🤔

- Batch gradient descent
  - ▶  $d$
- Stochastic gradient descent
  - ▶ 1
- Mini-batch gradient descent
  - ▶  $1 < b < d$

# Batching

	batch size	estimate	train time	computation
(Full) Batch gradient descent	$d$	accurate	high	high
Mini-Batch gradient descent	$1 < b < d$	less accurate	faster	faster
Stochastic gradient descent	1	least accurate	fastest	fasterst

## Notes

- All modes are (special cases of) mini-batch
- When people say “BGD” they usually mean “MBGD”

## Note on common memory saving technique

- Accumulate gradient for  $< b$  samples and update only after reaching mini-batch
- This way GPU/RAM only has  $< b$  samples at the time
- Output same as for minibatch
- No extra memory cost, higher computation time

# Batching




Figure 1: Convergence of SGD, MBGD, BGD [1]

# Optimization Algorithms - Momentum

## SGD/MBGD/BGD (known as SGD)

- Approximate gradient based on some set
- $\theta \leftarrow \theta + g$

## SGD with momentum

- Approximate gradient based on some set
- Update direction in which the gradient is moving (+exponential decay)
- $v \leftarrow \alpha v - \epsilon g$
- $\theta \leftarrow \theta + v$

## Remaining issue

- Some dimensions are very high frequency (need small learning rate)
- Some dimensions are uneventful (need large learning rate)

# Optimization Algorithms - Adaptive learning rates

## AdaGrad

- Keep sum of squared gradients (“speed”):  $r \leftarrow r + g \odot g$
- Scale individual dimensions:  $\theta \leftarrow \frac{\epsilon}{\delta + \sqrt{r}} \odot g$

## RMSProp

- Same as AdaGrad but decay history
- Keep sum of squared gradients (“speed”):  $r \leftarrow \rho r + (1 - \rho)g \odot g$
- Scale individual dimensions:  $\theta \leftarrow \frac{\epsilon}{\delta + \sqrt{r}} \odot g$

## Adam

- Combine momentum and adaptive learning rates
- $s \leftarrow \rho_1 s + (1 - \rho_1)g, r \leftarrow \rho_2 r + (1 - \rho_2)g \odot g$
- Bias correction (matters only at the beginning):  $\hat{s} = s / (1 - \rho_1^t), \hat{r} = r / (1 - \rho_2^t)$
- Update:  $\theta \leftarrow \frac{\epsilon}{\delta + \sqrt{\hat{r}}} \hat{s}$



# Optimization Algorithms

## Notes

- Look up Nesterov look-ahead momentum update 2
- Adam is the go-to algorithm nowadays
- Be able to write down pseudocode for the optimizers (in more detail than these slides)




Figure 2: Comparison of optimizers [3]

See nice animated gradient optimizer visualizations.

# Assignment 7

Any questions?

# Resources

- 1 [dragonnotes.org/DeepLearning/Optimization](https://dragonnotes.org/DeepLearning/Optimization) (very good notes!)
- 2 [mlfromscratch.com/optimizers-explained/#/](https://mlfromscratch.com/optimizers-explained/#/)
- 3 [towardsdatascience.com/adam-latest-trends-in-deep-learning-optimization-6be9a291375c](https://towardsdatascience.com/adam-latest-trends-in-deep-learning-optimization-6be9a291375c)
- 4 gradient visualization